RAMKUMAR ("RAM") BHARATHAN, MBA, CA

Audit & Technology-Risk Leader | Evaluating AI Judgment Under Uncertainty | Creator of ControlTrace

San Francisco, CA / Dallas, TX • U.S. Work Authorization (No Visa Sponsorship Required)
ControlTrace: https://github.com/ramtoo-cell/controltrace • **GitHub:** https://github.com/ramtoo-cell • **Chicago Booth Story:** chicagobooth.edu/why-booth/stories/from-a-small-village-to-an-executive-mba

SUMMARY

Technology-risk and audit leader with nearly two decades of reaching defensible conclusions from incomplete, sometimes misleading evidence, and deciding when to escalate. I lead IT audits covering cloud, cyber, and AI at a Fortune-10 global financial institution and am building **ControlTrace** (https://github.com/ramtoo-cell/controltrace), a benchmark that scores AI agents on evidence-faithful reasoning and calibrated escalation under codified audit-evidence standards (PCAOB AS 1105, ISA 500). Chicago Booth Executive MBA (Honors), Chartered Accountant (ICAI), Purdue MS in AI (in progress).

RESEARCH & EVALUATION

- **ControlTrace (pilot — https://github.com/ramtoo-cell/controltrace):** Six audit and compliance investigation cases with PCAOB AS 1105 / ISA 500 ground truth and named traps (*salience*, *accusation-is-not-evidence*, *calibrated uncertainty*), plus minimal pairs that flip between resolve and escalate on a single piece of corroborating evidence. Built through an agentic Claude workflow where I define the construct and trap taxonomy, review and correct each case and its ground truth, and verify the scorer. A blind Claude baseline reached the right outcome in **5 of 6** cases yet cited evidence it should have excluded in **5 of 6** cases, and in **1 of 6** (CT-001) confirmed a violation on an unsubstantiated competitor complaint after computing that a vendor's disclosure had lapsed >3 years earlier. *Limits: six cases, one model, one sample per case, ground truth not yet independently validated.*
- **Conceptual Reasoning & Calibration Suites (github.com/ramtoo-cell):** Built open-source evaluation packages including *frontier-reasoning-bench* (step-level invariant verification across economics, law, policy, philosophy, and cognitive science), *policy-reasoning-bench* (72 structured cases with 3-part rubrics), *reasoning-eval-harness*, *reasoning-paired-evals* (McNemar & paired bootstrap inference), and *reasoning-rubric-reliability*.
- **AI Assurance & Publications:** Authored a ~90-page control playbook for an AI-first company (including agentic-identity and CI/CD controls), open-source MCP tool-call governance proxies (*agentgate*, *agent-guardian*), and publications on **TD Commons** on AI and controls.

EXPERIENCE

VP, IT / Cloud & AI Audit — Fortune-10 Global Financial Institution (Dec 2025 – Present)

- Lead a team auditing cloud infrastructure, cybersecurity, and artificial intelligence, including enterprise AI and agent-hosting infrastructure; evaluate evidentiary sufficiency and escalation thresholds across high-stakes control environments.

VP, Technology Risk, SOX & Internal Audit — Texas Capital Bank

- Built cloud security governance, IT SOX assurance, and continuous control verification architectures from zero.

Associate Director, Technology Risk & Assurance — PwC (16 Years; India, Dubai, US)

- Served **Google, Meta, Microsoft, and Amazon**; wrote internal audit and IT SOX playbooks for **Google** and **Microsoft**; worked on the **Waymo** and **Tesla** external audits, applying PCAOB AS 1105 and ISA 500 evidentiary standards across complex technology environments.

EDUCATION & CREDENTIALS

- **The University of Chicago Booth School of Business** — Executive MBA (With Honors)
- **Purdue University** — M.S. in Artificial Intelligence (Fall 2026 – In Progress)
- **Institute of Chartered Accountants of India (ICAI)** — Chartered Accountant (CA)